



Training vs Inference — Pure Compute (Nano4 H200, batch 128, AMP)

Throughput



Peak GPU Memory



Pure-compute inference favours the tiny MT models by ~44x (197k vs 4.4k r/s). But end-to-end (with tokenisation + I/O, benchmark_summary.csv) MT 13-mer s1 drops to 1,066 r/s— below NT-v2's 1,221 — so its real bottleneck is overlapping-13-mer tokenisation over a 33M-row vocab, not the model. MT 13-mer training is also memory-bound (16.5 GB) from the dense embedding gradient.